

This demo showcases a Learning Agent concept using Ollama, a tool to run powerful LLMs (Large Language Models) locally on your computer. We'll install the tool and a specific model called Mistral, and then use it to simulate how a Learning Agent understands, adapts, and responds to new information.

Step-by-Step Process

◆ Step 1: Install Ollama

1. Go to the official website: <https://ollama.ai>
2. Click on Download and install the version for your OS (Windows, macOS, or Linux).
3. After installation, open the terminal or command prompt.

Step 2: Run Ollama

Once installed, Ollama runs as a background service. You can check if it's working by running:

- `ollama --help`
- `ollama --version` to confirm its installed

Step 3: Pull the Mistral Model

- `ollama pull mistral`

This will download the model locally (~4–8GB). Once done, it's ready to use — no internet required afterward

or you can download any model as you wish from the Ollama website model's section.

Step 4: Run the Mistral Model

- `ollama run mistral`
- The terminal will turn into a chat interface.

You can have interaction with it now:

You: What is reinforcement learning?

Mistral: Reinforcement learning is a type of machine learning where agents learn by...

.....

⚠ Note: This demonstration using Mistral via Ollama is not a pure Learning Agent in the classical sense — it doesn't retain memory of past user interactions across sessions.

However, it offers a simple and powerful introduction to Large Language Models (LLMs) and how they process, understand, and generate intelligent responses based on input — forming the foundation of learning agents in modern AI systems.